Lessons from the trenches

Improving response by being "data wrangling" amateurs in AWS

Swetha Balla, BSides Budapest, May 2021

Agenda

- Challenge with IR in AWS
- Logs
- Data wrangling
- Metrics
- Visualisation
- Questions

How many “things” do you need to improve incident response in AWS?

Just one! “Data wrangling”. (with a lot of caveats!)

Lot of caveats (logs, damn'ed logs!)

- ~~The security team want me to enable ALL logs!~~ Have I enabled key logs?
- ~~My SIEM costs are high!~~ How do I store these logs?
- ~~I need to transfer logs off-platform to create dashboards!~~ How can I visualise logs?

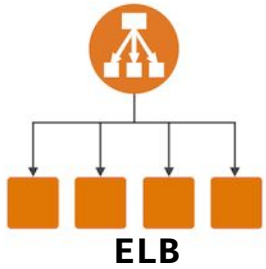

Data wrangling

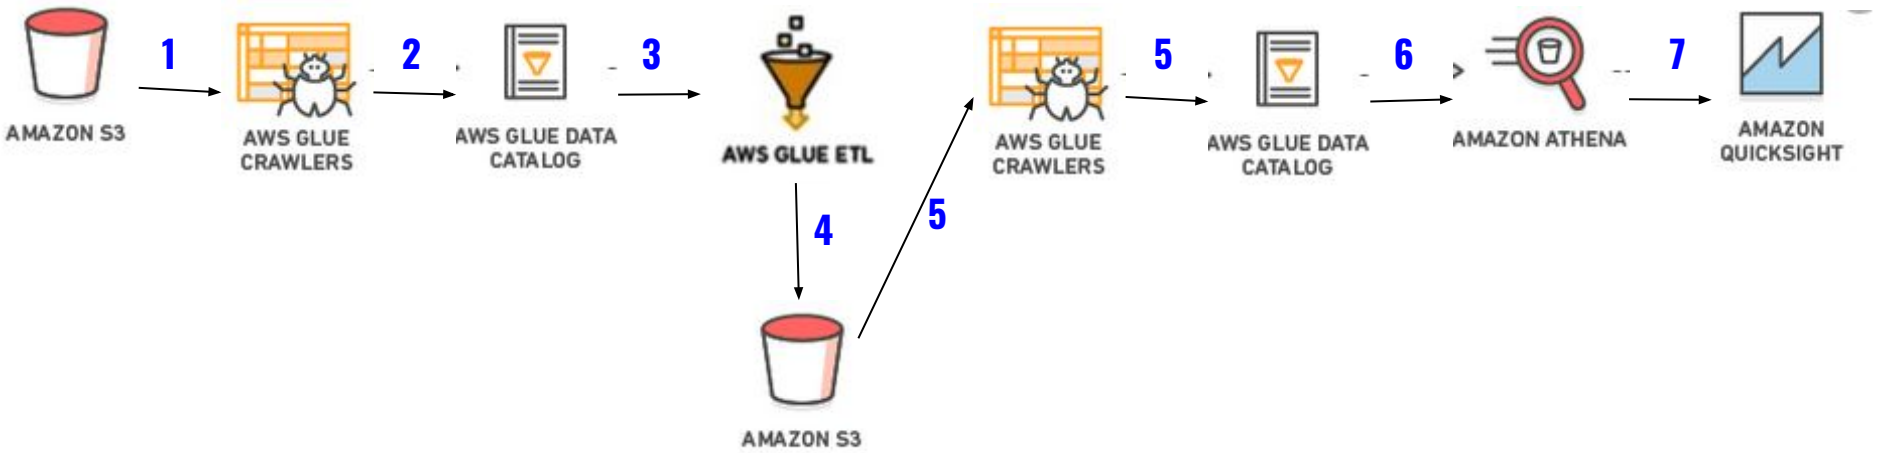

1. AWS Glue "Crawler" reads the CloudTrail logs from S3 (.json.gz)
2. Metadata for CloudTrail JSON logs added to AWS Glue "Tables". (**Note**: It's important to change the Data Type for the fields "*requestparameters*" and "*responseelements*" to string. By default, Glue sets them to struct)
3. AWS Glue ETL "jobs" convert JSON files to Parquet.
4. Converted parquet data stored in S3.
5. AWS Glue "Crawler" reads the CloudTrail logs from S3 (.json.gz).
6. Query parquet data in Athena
7. Visualise logs in Quicksight

Pre-requisites

- IAM role with the right permissions.
  - “AwsGlueServiceRole” policy
  - Other perms are also required - e.g. access to the S3 bucket with logs
- CloudTrail logs are available in S3 (.json.gz).
- Patience, in case something fails.

Demo

Example metrics

Query	json (unpartitioned)	parquet
API Errors SELECT eventTime, eventSource, eventName, errorCode, errorMessage, responseElements, awsRegion, userIdentity.arn, sourceIPAddress, userAgent FROM *<<table>>* WHERE errorCode IN ('Client.UnauthorizedOperation' , 'Client.InvalidPermission.NotFound ' , 'Client.OperationNotPermitted' ,'AccessDenied') ORDER BY eventTime DESC limit 25	Run time: 17.25 seconds Data scanned: 1.53 GB	Run time: 5.37 seconds Data scanned: 989.51 MB
Activity from malicious IP SELECT eventTime, eventSource, eventName, awsRegion, userIdentity.arn, sourceIPAddress, userAgent FROM *<<table>>* WHERE sourceIPAddress = '5.205.62.253' ORDER BY eventTime DESC limit 25	Run time: 19.32 seconds Data scanned: 1.53 GB	Run time: 2.98 seconds Data scanned: 9.05 MB
EC2 Instance enumerating S3 SELECT useridentity.principalid, eventsource, eventname, count(*) AS total FROM *<<table>>* WHERE useridentity.principalid LIKE '%:i-%' AND eventsource = 's3.amazonaws.com' AND eventname = 'ListBuckets' GROUP BY useridentity.principalid,eventsource,eventname ORDER BY total DESC limit 25	Run time: 13.23 seconds Data scanned: 1.53 GB	Run time: 1.19 seconds Data scanned: 38.87 MB

~ 74% less data scanned

~ 77% quicker

Parquet vs. JSON (unpartitioned) - query performance improvement

(admittedly, small data set!)

Example dashboard

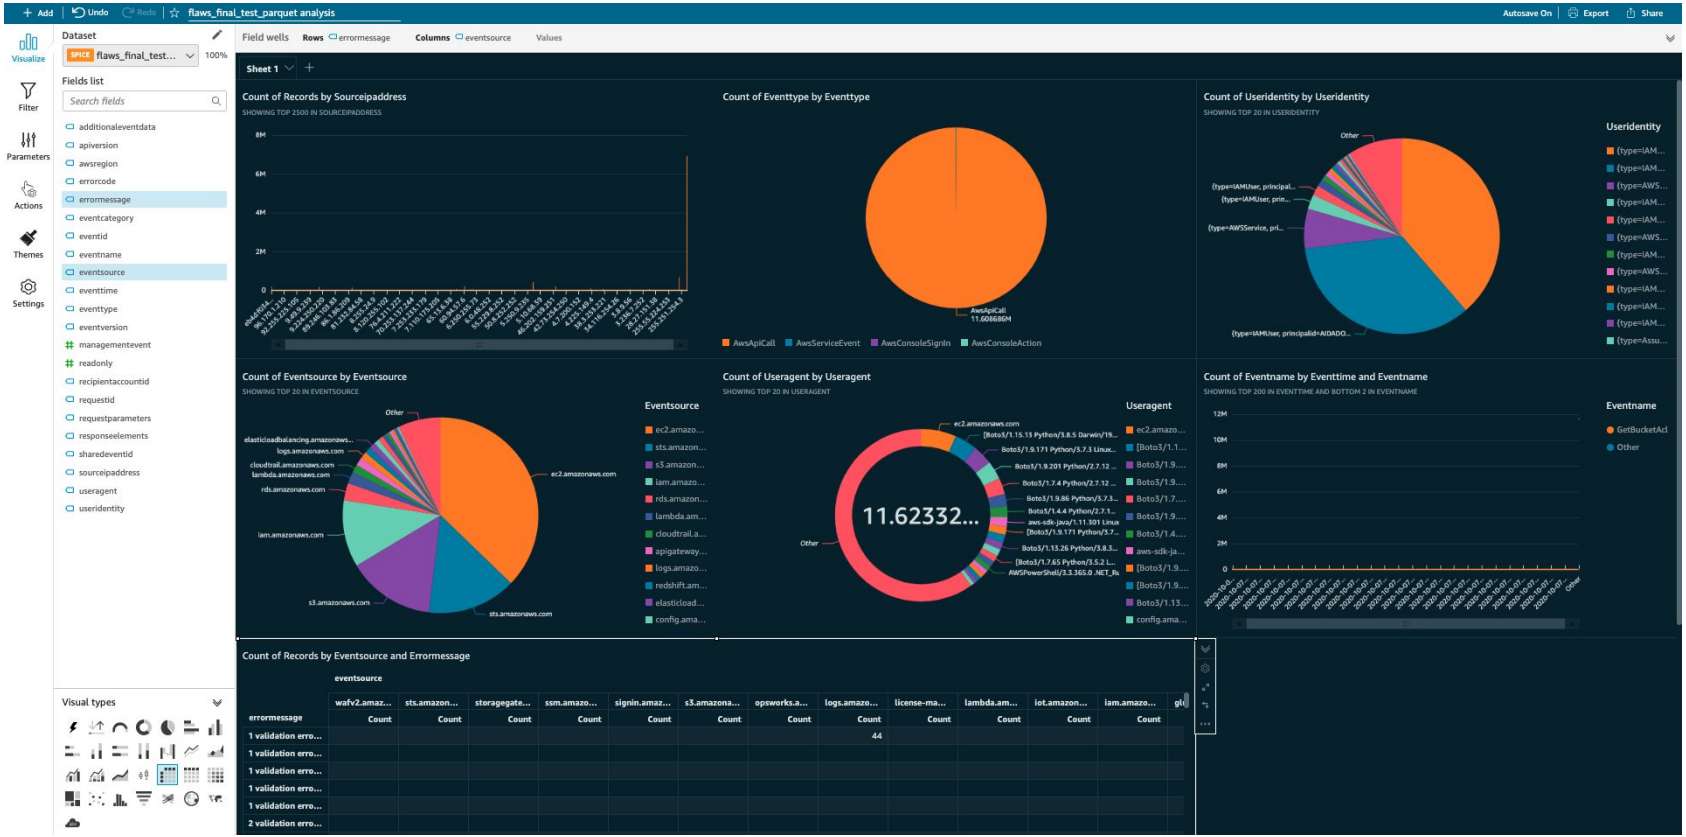

Take-aways

- ~~The security team want me to enable ALL logs!~~ Have I enabled key logs?
- ~~My SIEM costs are high!~~ How do I store these logs?
- ~~I need to transfer logs off-platform to create dashboards!~~ How can I visualise logs?

References

- Logging in the cloud: https://ponderthebits.com/wp-content/uploads/2020/02/Logging-in-the-Cloud-From-Zero-to-Incident-Response-Hero-Public.pdf
- Dataset: https://summitroute.com/blog/2020/10/09/public_dataset_of_cloudtrail_logs_from_flaws_cloud/ & http://summitroute.com/downloads/flaws_cloudtrail_logs.tar
- Example queries for AWS: https://github.com/easttimor/aws-incident-response
- AWS Glue:
  - https://aws.amazon.com/glue/
  - https://docs.aws.amazon.com/athena/latest/ug/glue-best-practices.html
  - https://aws.amazon.com/blogs/database/how-to-extract-transform-and-load-data-for-analytic-processing-using-aws-glue-part-2/
- AWS Athena: https://aws.amazon.com/athena/
- AWS Quicksight: https://aws.amazon.com/quicksight/
- Columnar data storage: https://docs.aws.amazon.com/athena/latest/ug/columnar-storage.html